

Homework 1

by

Krima Doshi

1222303329

CSE576: Natural Language Processing
Masters in Computer Science

Arizona State University
September 2, 2021

1 Using the Sentiment140 dataset available on Kaggle, apply BERT Encoding and MLP to train, validate and test a model for sentiment analysis.

1.1 Comparison of hyper-parameters tuned in model:

Batch Size	Max length	Layer Topology	Epoch	Learning Rate	Accuracy (Training, Validation, Testing)
64	512	(512, 256, 64, 32, 1)	(20,5)	0.5	80.48, 73.37, 73.20
64	512	(512, 256, 64, 32, 1)	25	0.5	80.48, 73.37, 73.20
64	512	(512, 256, 64, 32, 1)	25	0.1	82.67, 67.07, 67.10
64	512	(512, 256, 64, 32, 1)	50	0.1	89.34, 68.80, 66.23
64	512	(512, 256, 64, 32, 1)	100	0.01	74.62, 59.79, 60.29
64	64	(64, 128, 64, 32, 1)	47	0.5	81.49, 68.16, 68.07
64	64	(64, 256, 64, 32, 1)	50	0.1	70.01, 57.5, 57.91
16	32	(64, 256, 64, 32, 1)	50	0.1	80.27, 68.6, 67.98
16	64	(64, 256, 64, 32, 1)	(32,15)	(0.1,0.01)	80.95, 67.96, 67.67

1.2 Observations/Notes:

- Only one output node and translated them into 0s and 1s using a threshold 0.
- 5 epochs were not sufficient to get the desired accuracy so the model was trained for longer epochs.
- With a learning rate of 0.1, the model was taking too long to train and converge, so a higher learning rate was used. In some cases, the model also got stuck at the local minima with that learning rate.

- For some parameters, instead of training continuously, resetting the SGD helped increase the accuracy.

1.3 Output Graphs:

Graphs of some noteworthy observations:

1. I would consider this case the best as it not only got a good accuracy but predicted the given sentences correctly.

```
[53] # Post-training evaluation
evaluate_model(train_inputs, train_labels, data_name="Training")
evaluate_model(validation_inputs, validation_labels, data_name="Validation")
evaluate_model(test_inputs, test_labels, data_name="Testing")

Training
Loss: 0.6042267084121704
Accuracy: 0.8048285714285714
Validation
Loss: 0.6400972604751587
Accuracy: 0.7337333333333333
Testing
Loss: 0.6404624581336975
Accuracy: 0.7320333333333333
```

```
def predict_sentence(sentence):
    test_sents=[sentence]
    test_sents = [tokenizer.convert_tokens_to_ids(tokenize_truncate(sent, tokenizer, max_input_length)) for sent in test_sents]
    sent_inputs = pad_sequences(test_sents, maxlen=max_input_length, dtype="long", truncating="post", padding="post")
    sent_inputs = torch.LongTensor(sent_inputs)
    # test_labels = torch.FloatTensor(test_labels.to_list())
    # Initiate model in evaluation mode
    model.eval()
    # Squeeze output to 1D
    pred_labels = model(sent_inputs)
    return "pos" if (pred_labels>THRESHOLD).float() else "neg"

print(predict_sentence("It was a good experience"))
print(predict_sentence("It was a bad experience"))
print(predict_sentence("That was terrible"))
```

```
pos
neg
neg
```

Figure 1: Model Accuracy and Loss with 20 and 5 epochs and 0.5 learning rate.




Figure 2: Training Accuracy and validation accuracies per epoch

```

# Post-training evaluation
evaluate_model(train_inputs, train_labels, data_name="Training")
evaluate_model(validation_inputs, validation_labels, data_name="Validation")
evaluate_model(test_inputs, test_labels, data_name="Testing")

Training
Loss: 0.5899195671081543
Accuracy: 0.8267714285714286
Validation
Loss: 0.6674506664276123
Accuracy: 0.6706666666666666
Testing
Loss: 0.6665695905685425
Accuracy: 0.6710666666666667

def predict_sentence(sentence):
    test_sents=[sentence]
    test_sents = [tokenizer.convert_tokens_to_ids(tokenize_truncate(sent, tokenizer, max_input_length)) for sent in test_sents]
    sent_inputs = pad_sequences(test_sents, maxlen=max_input_length, dtype="long", truncating="post", padding="post")
    sent_inputs = torch.LongTensor(sent_inputs)
    # test_labels = torch.FloatTensor(test_labels.to_list())
    # Initiate model in evaluation mode
    model.eval()
    # Squeeze output to 1D
    pred_labels = model(sent_inputs)
    return "pos" if (pred_labels>THRESHOLD).float() else "neg"

print(predict_sentence("It was a good experience"))
print(predict_sentence("It was a bad experience"))
print(predict_sentence("That was terrible"))

pos
neg
pos

```

Figure 3: Model Accuracy and Loss with 25 epochs and 0.1 learning rate.



Figure 4: Training Accuracy and validation accuracies per epoch

```
[18] # Post-training evaluation
evaluate_model(train_inputs, train_labels, data_name="Training")
evaluate_model(validation_inputs, validation_labels, data_name="Validation")
evaluate_model(test_inputs, test_labels, data_name="Testing")

Training
  Loss: 0.5562483072280884
  Accuracy: 0.8934
Validation
  Loss: 0.6682621836662292
  Accuracy: 0.6680666666666667
Testing
  Loss: 0.6710196733474731
  Accuracy: 0.6623666666666667
```

```
def predict_sentence(sentence):
    test_sents=[sentence]
    test_sents = [tokenizer.convert_tokens_to_ids(tokenize_truncate(sent, tokenizer, max_input_length)) for sent in test_sents]
    sent_inputs = pad_sequences(test_sents, maxlen=max_input_length, dtype="long", truncating="post", padding="post")
    sent_inputs = torch.LongTensor(sent_inputs)
    # test_labels = torch.FloatTensor(test_labels.to_list())
    # Initiate model in evaluation mode
    model.eval()
    # Squeeze output to 1D
    pred_labels = model(sent_inputs)
    return "pos" if (pred_labels>THRESHOLD).float() else "neg"

print(predict_sentence("It was a good experience"))
print(predict_sentence("It was a bad experience"))
print(predict_sentence("That was terrible")) |
```

```
pos
neg
pos
```

Figure 5: Model Accuracy and Loss with 50 epochs and 0.1 learning rate.

3.

```
[25] # Post-training evaluation
evaluate_model(train_inputs, train_labels, data_name="Training")
evaluate_model(validation_inputs, validation_labels, data_name="Validation")
evaluate_model(test_inputs, test_labels, data_name="Testing")

Training
  Loss: 0.6222451329231262
  Accuracy: 0.7462857142857143
Validation
  Loss: 0.6969066858291626
  Accuracy: 0.5979333333333333
Testing
  Loss: 0.6943069100379944
  Accuracy: 0.6029
(0.6943069100379944, 0.6029)
```

```
def predict_sentence(sentence):
    test_sents=[sentence]
    test_sents = [tokenizer.convert_tokens_to_ids(tokenize_truncate(sent, tokenizer, max_input_length)) for sent in test_sents]
    sent_inputs = pad_sequences(test_sents, maxlen=max_input_length, dtype="long", truncating="post", padding="post")
    sent_inputs = torch.LongTensor(sent_inputs)
    # test_labels = torch.FloatTensor(test_labels.to_list())
    # Initiate model in evaluation mode
    model.eval()
    # Squeeze output to 1D
    pred_labels = model(sent_inputs)
    return "pos" if (pred_labels>THRESHOLD).float() else "neg"

print(predict_sentence("It was a good experience"))
print(predict_sentence("It was a bad experience"))
print(predict_sentence("That was terrible"))
```

```
pos
pos
neg
```

Figure 6: Model Accuracy and Loss with 100 epochs and 0.01 learning rate.

4.